Chapter 3

Instruction-Level Parallelism and Its Exploitation

Introduction

- Pipelining become universal technique in 1985
  - Overlaps execution of instructions
  - Exploits "Instruction Level Parallelism"

- Beyond this, there are two main approaches:
  - Hardware-based dynamic approaches
    - Used in server and desktop processors
    - Not used as extensively in PMP processors
  - Compiler-based static approaches
    - Not as successful outside of scientific applications

Instruction-Level Parallelism

- When exploiting instruction-level parallelism, goal is to maximize CPI
  - Pipeline CPI =
    - Ideal pipeline CPI +
    - Structural stalls +
    - Data hazard stalls +
    - Control stalls

- Parallelism with basic block is limited
  - Typical size of basic block = 3-6 instructions
  - Must optimize across branches

Data Dependence

- ## Loop-Level Parallelism
  - Unroll loop statically or dynamically
  - Use SIMD (vector processors and GPUs)

- ## Challenges:
  - Data dependency
    - Instruction j is data dependent on instruction i if
      - Instruction i produces a result that may be used by instruction j
      - Instruction j is data dependent on instruction k and instruction k is data dependent on instruction i

- ## Dependent instructions cannot be executed simultaneously

Data Dependence

- Dependencies are a property of programs

- Pipeline organization determines if dependence is detected and if it causes a stall

- Data dependence conveys:
    - Possibility of a hazard
    - Order in which results must be calculated
    - Upper bound on exploitable instruction level parallelism

- Dependencies that flow through memory locations are difficult to detect

5

Name Dependence

- Two instructions use the same name but no flow of information

    - Not a true data dependence, but is a problem when reordering instructions

    - Antidependence: instruction j writes a register or memory location that instruction i reads

        - Initial ordering (i before j) must be preserved

    - Output dependence: instruction i and instruction j write the same register or memory location

        - Ordering must be preserved

- To resolve, use register renaming techniques

Other Factors

- ## Data Hazards
    - ### Read after write (RAW)
    - ### Write after write (WAW)
    - ### Write after read (WAR)

- ## Control Dependence
    - ### Ordering of instruction i with respect to a branch instruction
        - Instruction control dependent on a branch cannot be moved before the branch so that its execution is no longer controlled by the branch
        - An instruction not control dependent on a branch cannot be moved after the branch so that its execution is controlled by the branch

Examples

- Example 1:

  add x1,x2,x3

  beq x4,x0,L

  sub x1,x1,x6

  L: …

  or x7,x1,x8

- Example 2:

  add x1,x2,x3

  beq x12,x0,skip

  sub x4,x5,x6

  add x5,x4,x9

  skip:

  or x7,x8,x9

■ or instruction dependent on add and sub

■ Assume x4 isn't used after skip

  ■ Possible to move sub before the branch

MORGAN KAUFMANN

Compiler Techniques for Exposing ILP

- ## Pipeline scheduling
  - Separate dependent instruction from the source instruction by the pipeline latency of the source instruction
- ## Example:

  for (i=999; i>=0; i=i-1)

     x[i] = x[i] + s;

Instruction producing result	Instruction using result	Latency in clock cycles
FP ALU op	Another FP ALU op	3
FP ALU op	Store double	2
Load double	FP ALU op	1
Load double	Store double	0

Pipeline Stalls

Loop: fld   f0,0(x1)
         stall
         fadd.d f4,f0,f2
         stall
         stall
         fsd f4,0(x1)
         addi x1,x1,-8
         bne x1,x2,Loop

Loop: fld   f0,0(x1)
         addi x1,x1,-8
         fadd.d f4,f0,f2
         stall
         stall
         fsd f4,0(x1)
         bne x1,x2,Loop

Instruction producing result	Instruction using result	Latency in clock cycles
FP ALU op	Another FP ALU op	3
FP ALU op	Store double	2
Load double	FP ALU op	1
Load double	Store double	0

Loop Unrolling

- Loop unrolling
    - Unroll by a factor of 4 (assume # elements is divisible by 4)
    - Eliminate unnecessary instructions

```
Loop:      fld f0,0(x1)
           fadd.d f4,f0,f2
           fsd f4,0(x1) //drop addi & bne
           fld f6,-8(x1)
           fadd.d f8,f6,f2
           fsd f8,-8(x1) //drop addi & bne
           fld f0,-16(x1)
           fadd.d f12,f0,f2
           fsd f12,-16(x1) //drop addi & bne
           fld f14,-24(x1)
           fadd.d f16,f14,f2
           fsd f16,-24(x1)
           addi x1,x1,-32
           bne x1,x2,Loop
```

- note: number of live registers vs. original loop

Loop Unrolling/Pipeline Scheduling

- Pipeline schedule the unrolled loop:

```
Loop:      fld f0,0(x1)
           fld f6,-8(x1)
           fld f8,-16(x1)
           fld f14,-24(x1)
           fadd.d f4,f0,f2
           fadd.d f8,f6,f2
           fadd.d f12,f0,f2
           fadd.d f16,f14,f2
           fsd f4,0(x1)
           fsd f8,-8(x1)
           fsd f12,-16(x1)
           fsd f16,-24(x1)
           addi x1,x1,-32
           bne x1,x2,Loop
```

- 14 cycles
- 3.5 cycles per element

Strip Mining

- Unknown number of loop iterations?
  - Number of iterations = n
  - Goal: make k copies of the loop body
  - Generate pair of loops:
    - First executes n mod k times
    - Second executes n / k times
   - "Strip mining"

Branch Prediction

- ## Basic 2-bit predictor:
  - For each branch:
    - Predict taken or not taken
    - If the prediction is wrong two consecutive times, change prediction
- ## Correlating predictor:
  - Multiple 2-bit predictors for each branch
  - One for each possible combination of outcomes of preceding n branches
    - (m,n) predictor: behavior from last m branches to choose from 2^m n-bit predictors
- ## Tournament predictor:
 - Combine correlating predictor with local predictor

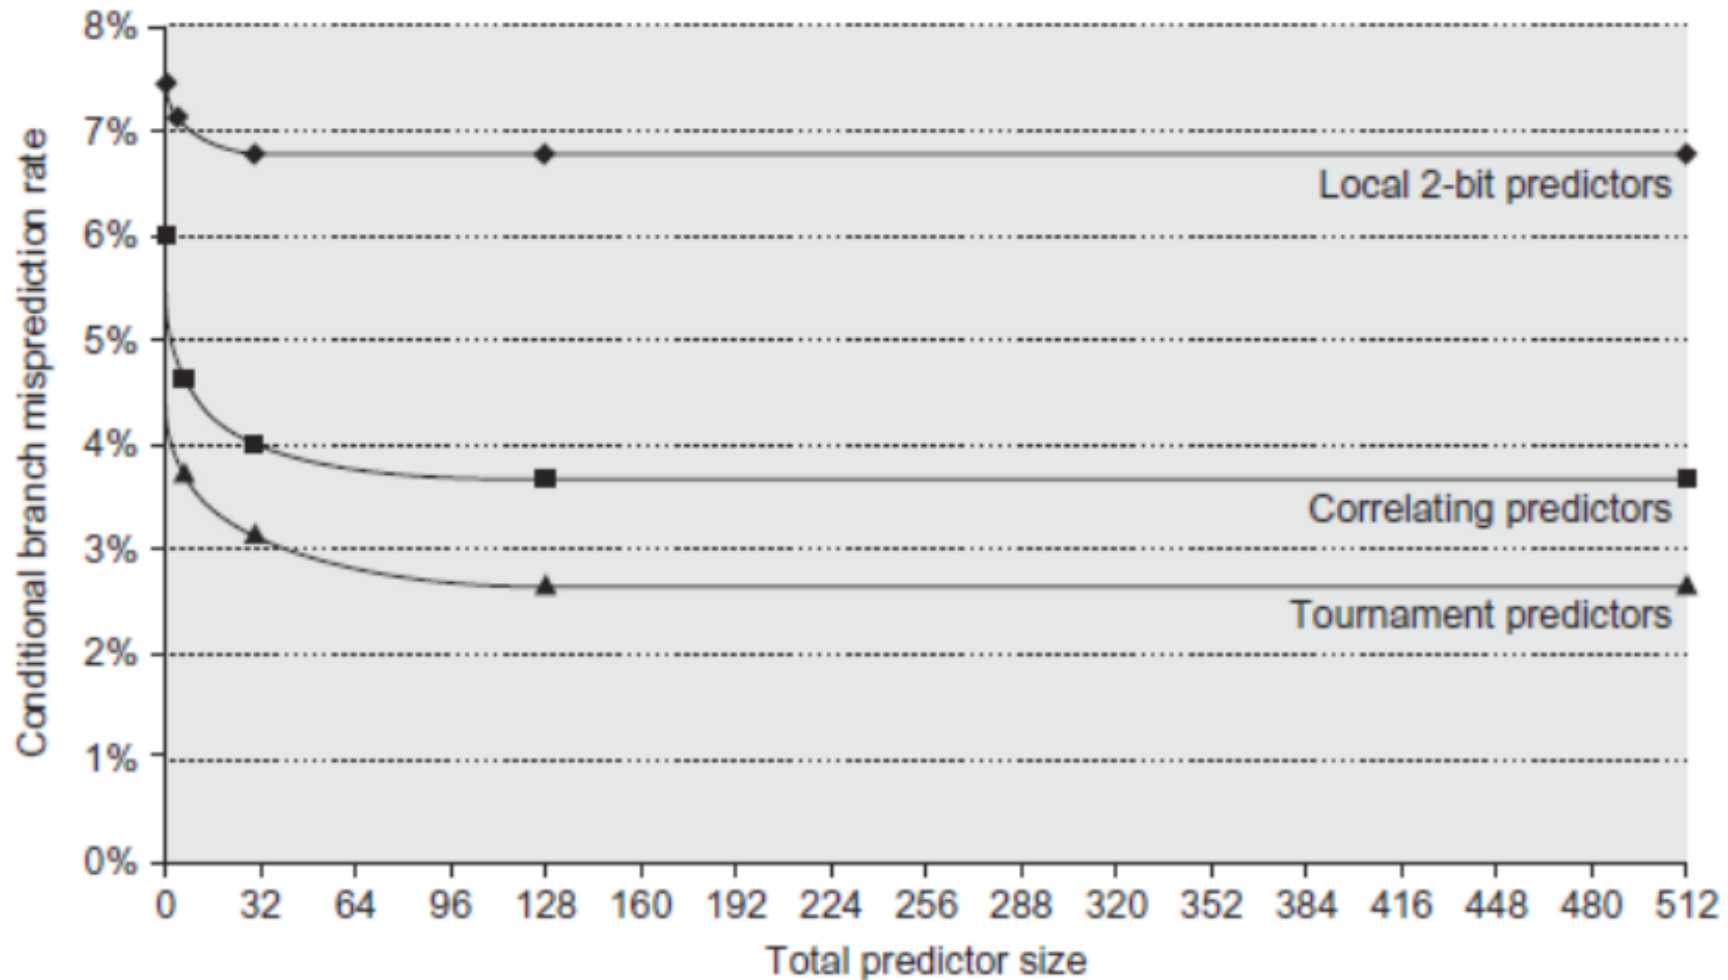# Branch Prediction
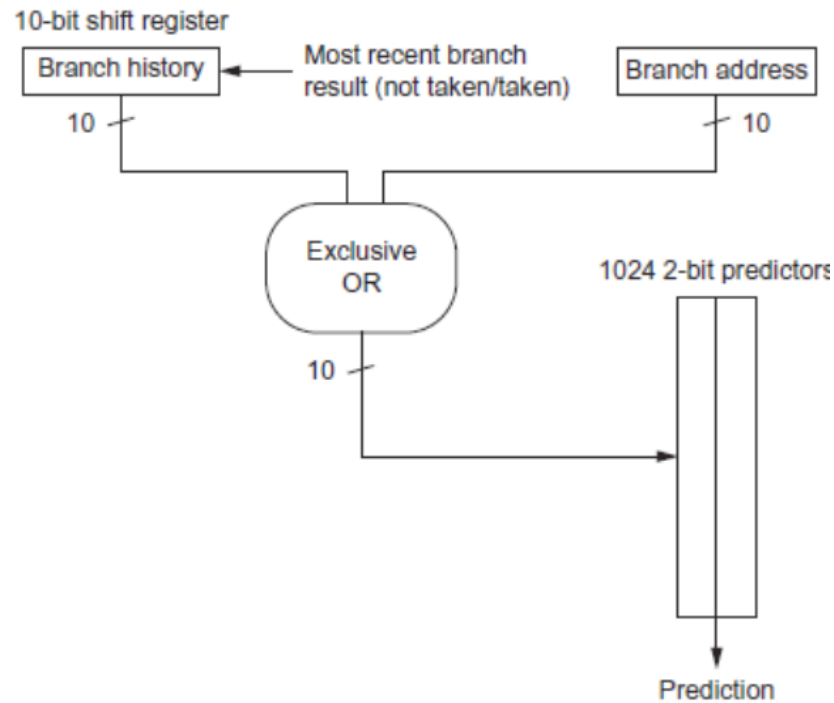
gshare
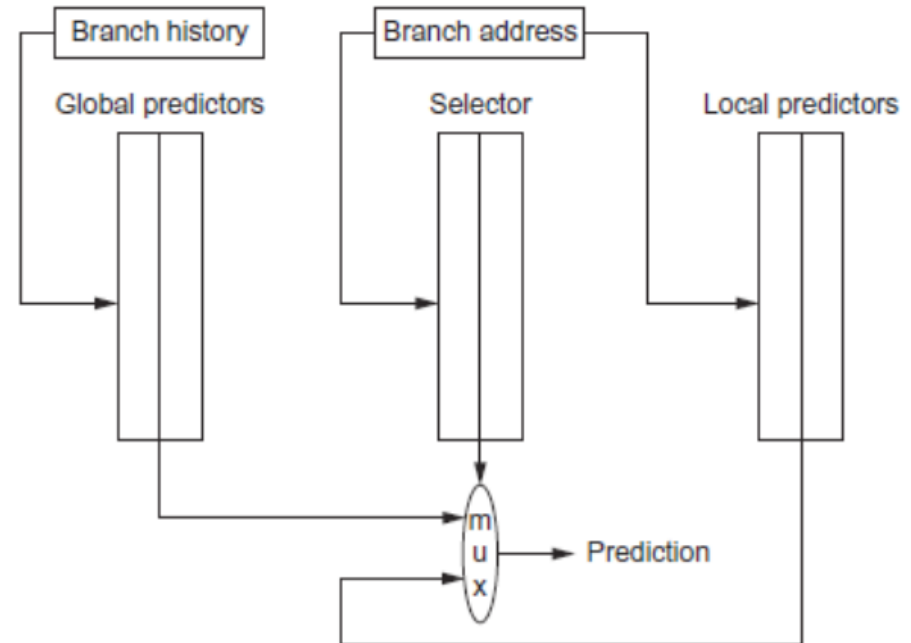
tournament

Branch Prediction Performance

Branch Prediction Performance

Tagged Hybrid Predictors

- Need to have predictor for each branch and history

  - Problem: this implies huge tables

  - Solution:

    - Use hash tables, whose hash value is based on branch address and branch history

    - Longer histories may lead to increased chance of hash collision, so use multiple tables with increasingly shorter histories

Tagged Hybrid Predictors

19

Tagged Hybrid Predictors

20
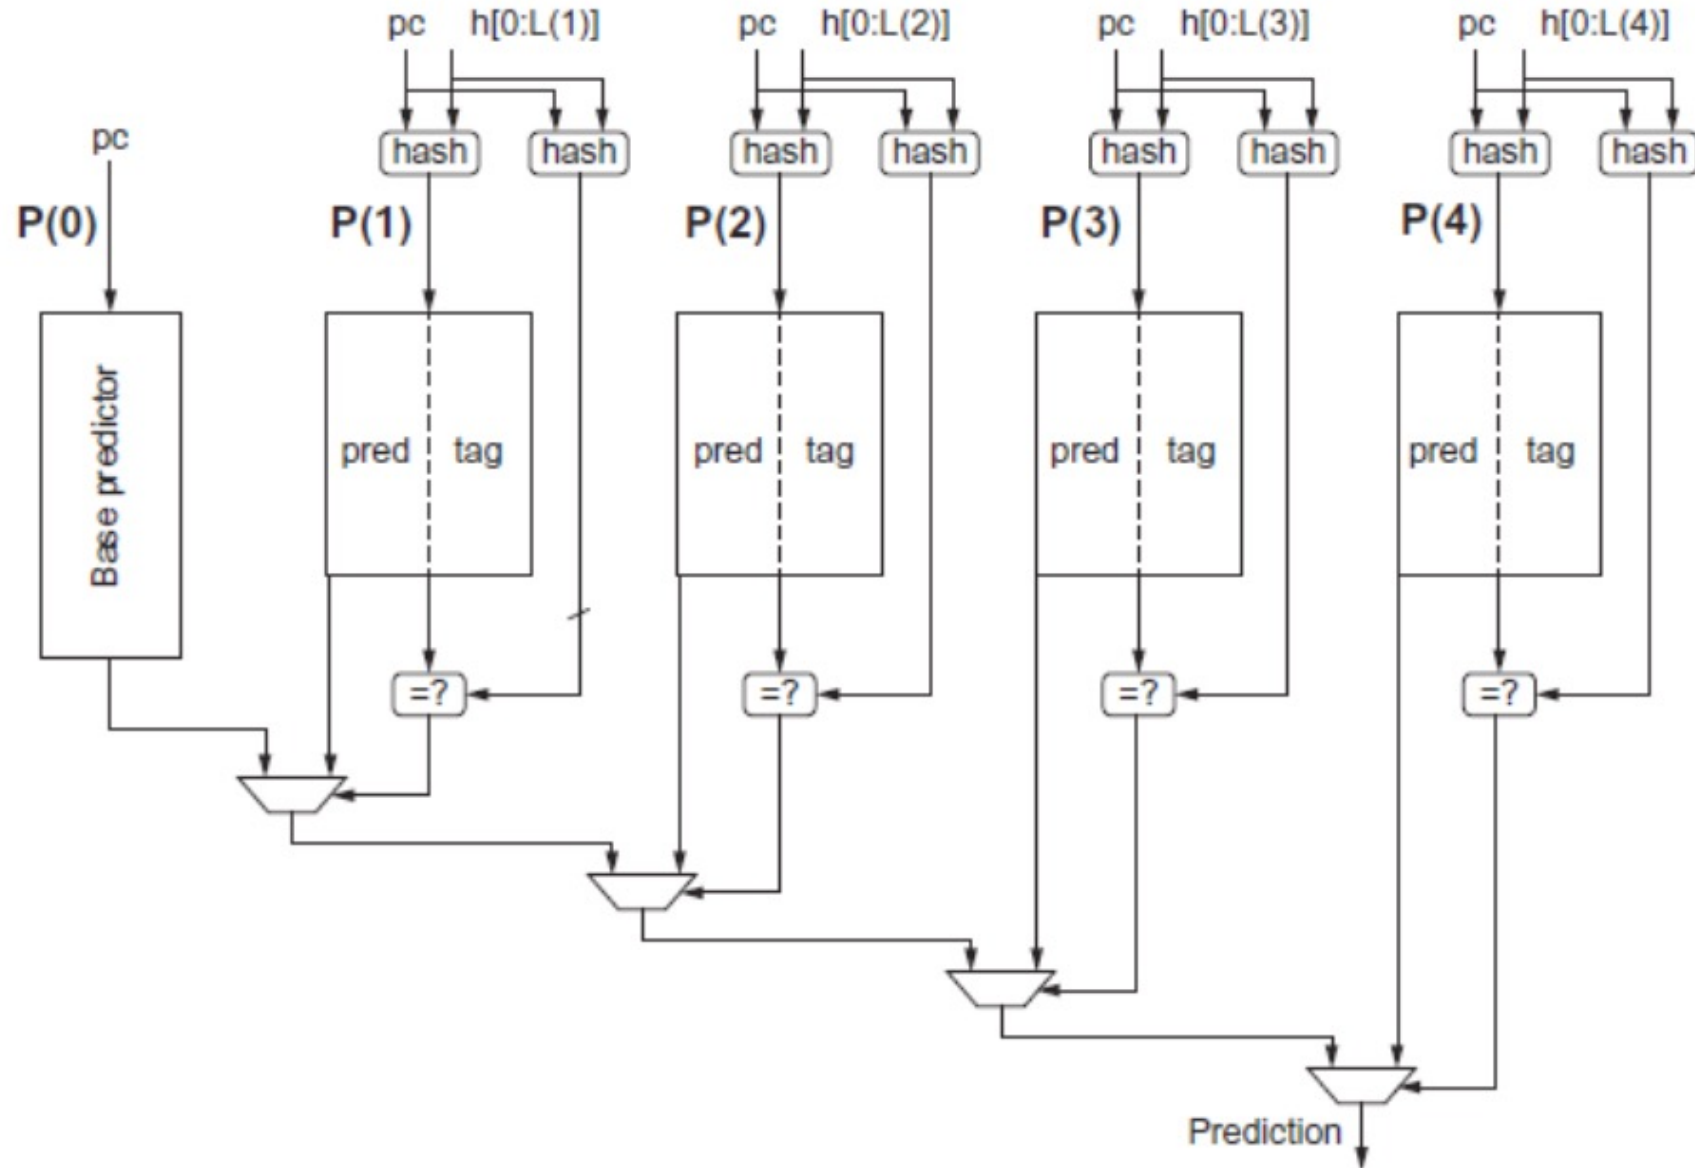
Dynamic Scheduling

- Rearrange order of instructions to reduce stalls while maintaining data flow

- Advantages:
  - Compiler doesn't need to have knowledge of microarchitecture
  - Handles cases where dependencies are unknown at compile time

- Disadvantage:
  - Substantial increase in hardware complexity
  - Complicates exceptions

Dynamic Scheduling

- Dynamic scheduling implies:
  - Out-of-order execution
  - Out-of-order completion

- Example 1:

  fdiv.d f0,f2,f4

  fadd.d f10,f0,f8

  fsub.d f12,f8,f14

  - fsub.d is not dependent, issue before fadd.d

Dynamic Scheduling

- Example 2:

  fdiv.d f0,f2,f4

  fmul.d f6,f0,f8

  fadd.d f0,f10,f14

  - fadd.d is not dependent, but the antidependence makes it impossible to issue earlier without register renaming

Register Renaming

- Example 3:

fdiv.d f0,f2,f4

fadd.d f6,f0,f8

fsd f6,0(x1)

fsub.d f8,f10,f14

fmul.d f6,f10,f8

antidependence

antidependence

- name dependence with f6

Register Renaming

- Example 3:

fdiv.d f0,f2,f4

fadd.d S,f0,f8

fsd S,0(x1)

fsub.d T,f10,f14

fmul.d f6,f10,T

- Now only RAW hazards remain, which can be strictly ordered

Register Renaming

- ## Tomasulo's Approach

  - ### Tracks when operands are available

  - ### Introduces register renaming in hardware

    - Minimizes WAW and WAR hazards

- ## Register renaming is provided by reservation stations (RS)

  - ### Contains:

    - The instruction

    - Buffered operand values (when available)

    - Reservation station number of instruction providing the operand values

Register Renaming

- RS fetches and buffers an operand as soon as it becomes available (not necessarily involving register file)
- Pending instructions designate the RS to which they will send their output
    - Result values broadcast on a result bus, called the common data bus (CDB)
- Only the last output updates the register file
- As instructions are issued, the register specifiers are renamed with the reservation station
- May be more reservation stations than registers
- Load and store buffers
    - Contain data and addresses, act like reservation stations

27

Tomasulo's Algorithm
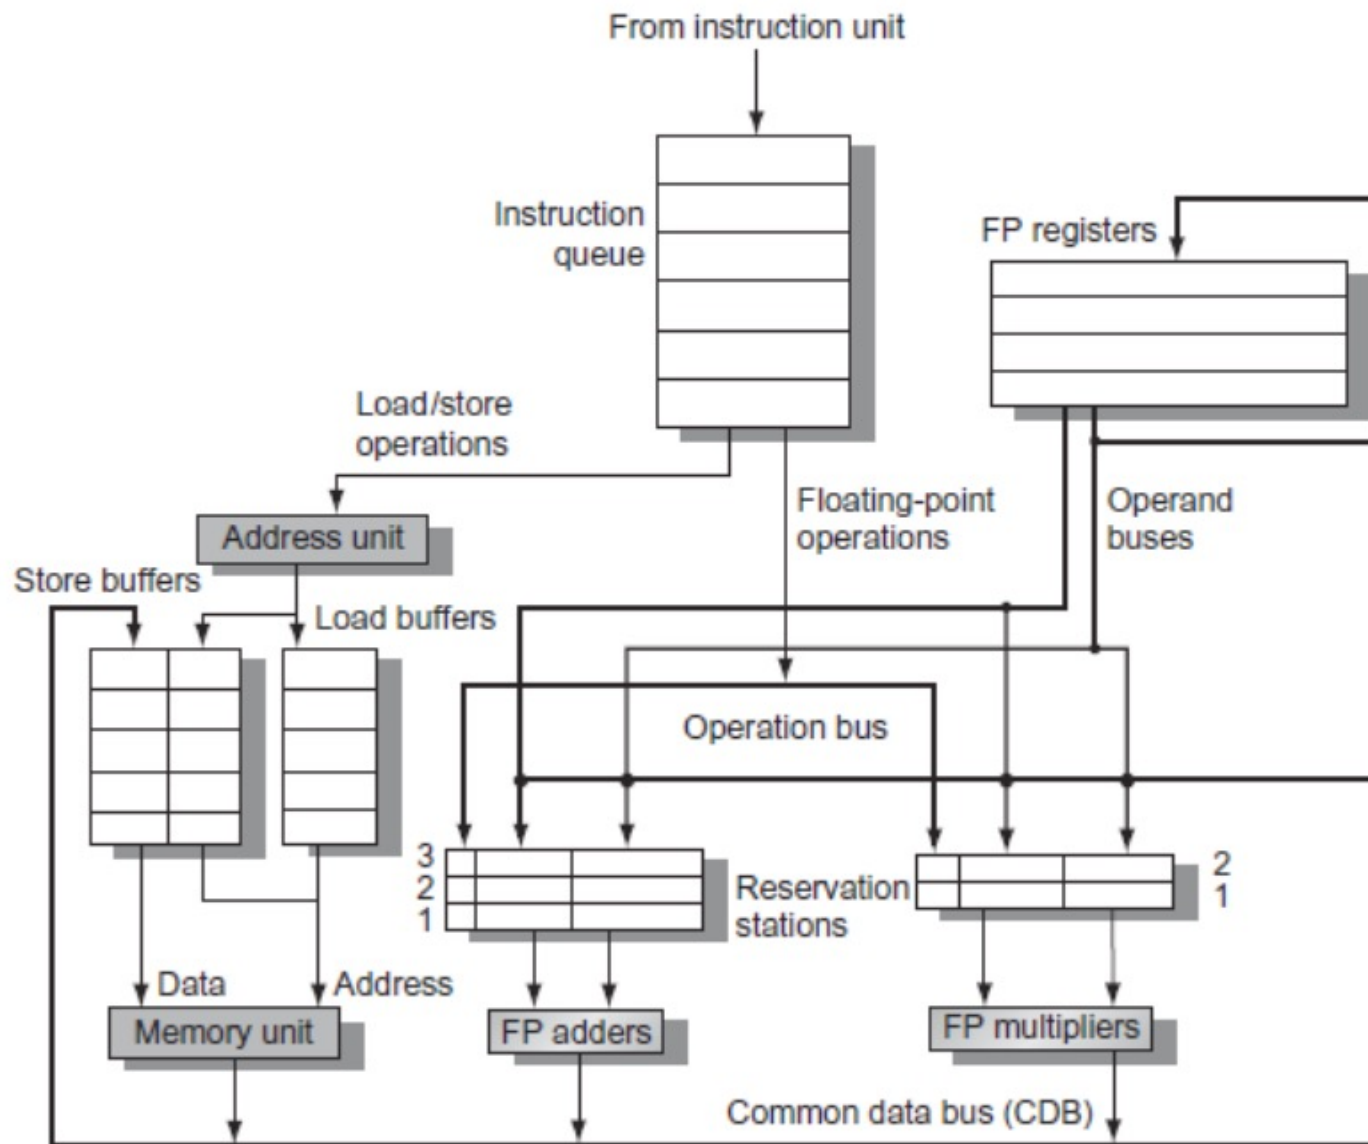
Tomasulo's Algorithm

- Three Steps:
  - Issue
    - Get next instruction from FIFO queue
    - If available RS, issue the instruction to the RS with operand values if available
    - If operand values not available, stall the instruction
  - Execute
    - When operand becomes available, store it in any reservation stations waiting for it
    - When all operands are ready, issue the instruction
    - Loads and store maintained in program order through effective address
    - No instruction allowed to initiate execution until all branches that proceed it in program order have completed
  - Write result
    - Write result on CDB into reservation stations and store buffers
      - (Stores must wait until address and value are received)

Example

Instruction status

Instruction		Issue	Execute	Write result
fld	f6,32(x2)	√	√	√
fld	f2,44(x3)	√	√	
fmul.d	f0,f2,f4	√		
fsub.d	f8,f2,f6	√		
fdiv.d	f0,f0,f6	√		
fadd.d	f6,f8,f2	√		

Reservation stations

Name	Busy	Op	Vj	Vk	Qj	Qk	A
Load1	No						
Load2	Yes	Load					44 + Regs[x3]
Add1	Yes	SUB		Mem[32 + Regs[x2]]	Load2		
Add2	Yes	ADD			Add1	Load2	
Add3	No						
Mult1	Yes	MUL		Regs[f4]	Load2		
Mult2	Yes	DIV		Mem[32 + Regs[x2]]	Mult1		

Register status

Field	f0	f2	f4	f6	f8	f10	f12	...	f30
Qi	Mult1	Load2		Add2	Add1	Mult2			

Tomasulo's Algorithm

- Example loop:

```
Loop:    fld f0,0(x1)
         fmul.d f4,f0,f2
         fsd f4,0(x1)
         addi x1,x1,8
         bne x1,x2,Loop // branches if x16 != x2
```

Tomasulo's Algorithm

Instruction		Instruction status			
		From iteration	Issue	Execute	Write result
fld	f0,0(x1)	1	√	√	
fmul.d	f4,f0,f2	1	√		
fsd	f4,0(x1)	1	√		
fld	f0,0(x1)	2	√	√	
fmul.d	f4,f0,f2	2	√		
fsd	f4,0(x1)	2	√		

				Reservation stations				
Name	Busy	Op	Vj	Vk	Qj	Qk	A	
Load1	Yes	Load					Regs[x1] + 0	
Load2	Yes	Load					Regs[x1] − 8	
Add1	No							
Add2	No							
Add3	No							
Mult1	Yes	MUL		Regs[f2]	Load1			
Mult2	Yes	MUL		Regs[f2]	Load2			
Store1	Yes	Store	Regs[x1]			Mult1		
Store2	Yes	Store	Regs[x1] − 8			Mult2		

				Register status					
Field	f0	f2	f4	f6	f8	f10	f12	...	f30
Qi	Load2		Mult2						

Hardware-Based Speculation

- Execute instructions along predicted execution paths but only commit the results if prediction was correct

- Instruction commit: allowing an instruction to update the register file when instruction is no longer speculative

- Need an additional piece of hardware to prevent any irrevocable action until an instruction commits

  - I.e. updating state or taking an execution

Reorder Buffer

- Reorder buffer – holds the result of instruction between completion and commit

- Four fields:
  - Instruction type: branch/store/register
  - Destination field: register number
  - Value field: output value
  - Ready field: completed execution?

- Modify reservation stations:
 - Operand source is now reorder buffer instead of functional unit

Reorder Buffer

- Issue:
  - Allocate RS and ROB, read available operands
- Execute:
  - Begin execution when operand values are available
- Write result:
  - Write result and ROB tag on CDB
- Commit:
  - When ROB reaches head of ROB, update register
  - When a mispredicted branch reaches head of ROB, discard all entries

Reorder Buffer

- Register values and memory values are not written until an instruction commits

- On misprediction:
  - Speculated entries in ROB are cleared

- Exceptions:
  - Not recognized until it is ready to commit

36
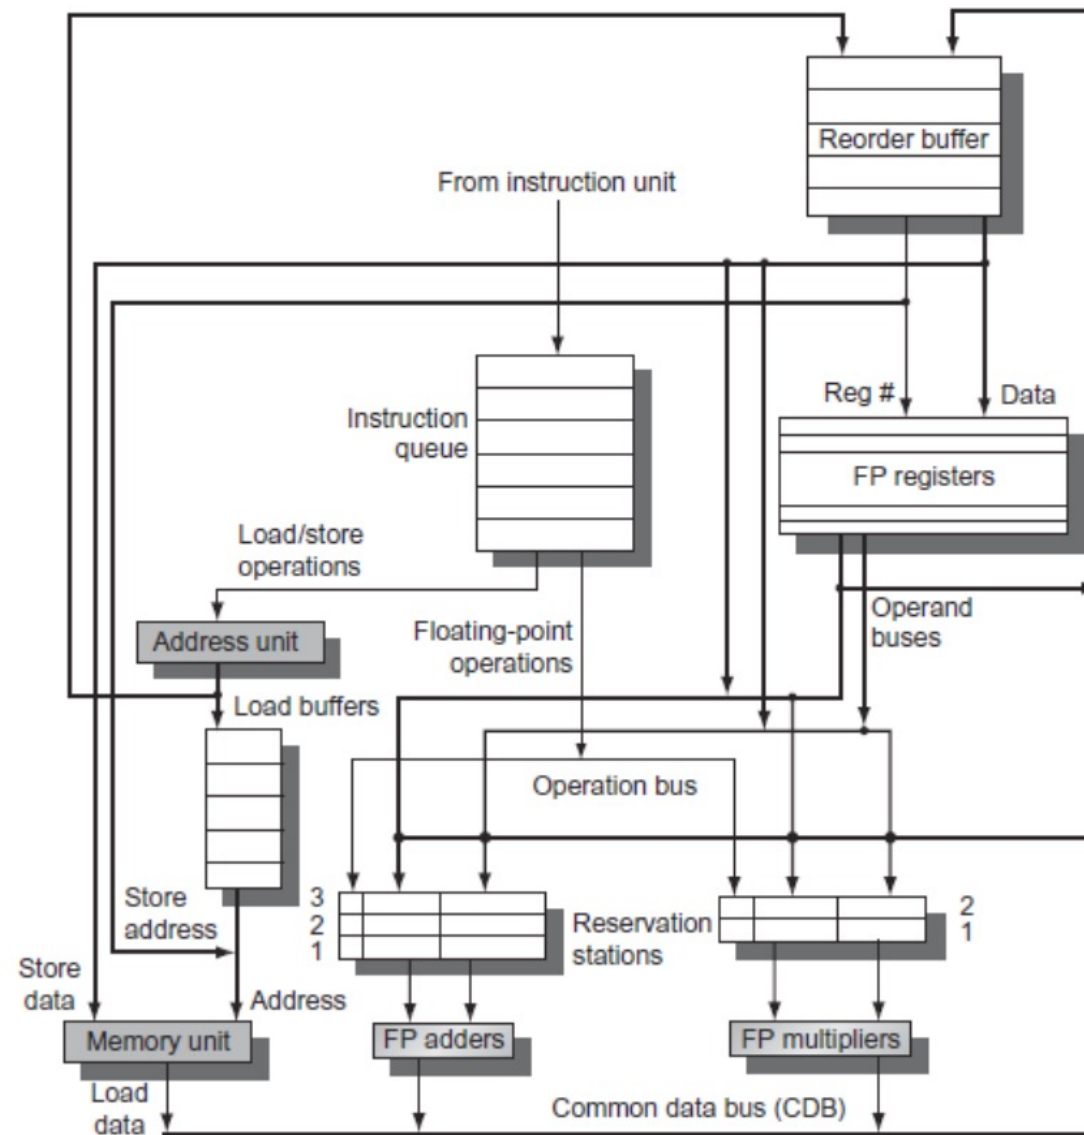
Reorder Buffer

Reorder Buffer

Reorder buffer

Entry	Busy	Instruction		State	Destination	Value
1	No	fld	f6,32(x2)	Commit	f6	Mem[32 + Regs[x2]]
2	No	fld	f2,44(x3)	Commit	f2	Mem[44 + Regs[x3]]
3	Yes	fmul.d	f0,f2,f4	Write result	f0	#2 × Regs[f4]
4	Yes	fsub.d	f8,f2,f6	Write result	f8	#2 − #1
5	Yes	fdiv.d	f0,f0,f6	Execute	f0	
6	Yes	fadd.d	f6,f8,f2	Write result	f6	#4 + #2

Reservation stations

Name	Busy	Op	Vj	Vk	Qj	Qk	Dest	A
Load1	No							
Load2	No							
Add1	No							
Add2	No							
Add3	No							
Mult1	No	fmul.d	Mem[44 + Regs[x3]]	Regs[f4]			#3	
Mult2	Yes	fdiv.d		Mem[32 + Regs[x2]]	#3		#5	

FP register status

Field	f0	f1	f2	f3	f4	f5	f6	f7	f8	f10
Reorder #	3						6		4	5
Busy	Yes	No	No	No	No	No	Yes	...	Yes	Yes

Multiple Issue and Static Scheduling

- To achieve CPI < 1, need to complete multiple instructions per clock

- Solutions:
    - Statically scheduled superscalar processors
    - VLIW (very long instruction word) processors
    - Dynamically scheduled superscalar processors

Multiple Issue

Common name	Issue structure	Hazard detection	Scheduling	Distinguishing characteristic	Examples
Superscalar (static)	Dynamic	Hardware	Static	In-order execution	Mostly in the embedded space: MIPS and ARM, including the Cortex-A53
Superscalar (dynamic)	Dynamic	Hardware	Dynamic	Some out-of-order execution, but no speculation	None at the present
Superscalar (speculative)	Dynamic	Hardware	Dynamic with speculation	Out-of-order execution with speculation	Intel Core i3, i5, i7; AMD Phenom; IBM Power 7
VLIW/LIW	Static	Primarily software	Static	All hazards determined and indicated by compiler (often implicitly)	Most examples are in signal processing, such as the TI C6x
EPIC	Primarily static	Primarily software	Mostly static	All hazards determined and indicated explicitly by the compiler	Itanium

VLIW Processors

- Package multiple operations into one instruction

- Example VLIW processor:
  - One integer instruction (or branch)
  - Two independent floating-point operations
  - Two independent memory references

- Must be enough parallelism in code to fill the available slots

VLIW Processors

Memory reference 1	Memory reference 2	FP operation 1	FP operation 2	Integer operation/branch
fld f0,0(x1)	fld f6,-8(x1)			
fld f10,-16(x1)	fld f14,-24(x1)			
fld f18,-32(x1)	fld f22,-40(x1)	fadd.d f4,f0,f2	fadd.d f8,f6,f2	
fld f26,-48(x1)		fadd.d f12,f0,f2	fadd.d f16,f14,f2	
		fadd.d f20,f18,f2	fadd.d f24,f22,f2	
fsd f4,0(x1)	fsd f8,-8(x1)	fadd.d f28,f26,f24		
fsd f12,-16(x1)	fsd f16,-24(x1)			addi x1,x1,-56
fsd f20,24(x1)	fsd f24,16(x1)			
fsd f28,8(x1)				bne x1,x2,Loop

- **Disadvantages:**
  - Statically finding parallelism
  - Code size
  - No hazard detection hardware
  - Binary code compatibility

Dynamic Scheduling, Multiple Issue, and Speculation

- ## Modern microarchitectures:

  - ### Dynamic scheduling + multiple issue + speculation

- ## Two approaches:

  - ### Assign reservation stations and update pipeline control table in half clock cycles

    - Only supports 2 instructions/clock

  - ### Design logic to handle any possible dependencies between the instructions

- ## Issue logic is the bottleneck in dynamically scheduled superscalars

Overview of Design
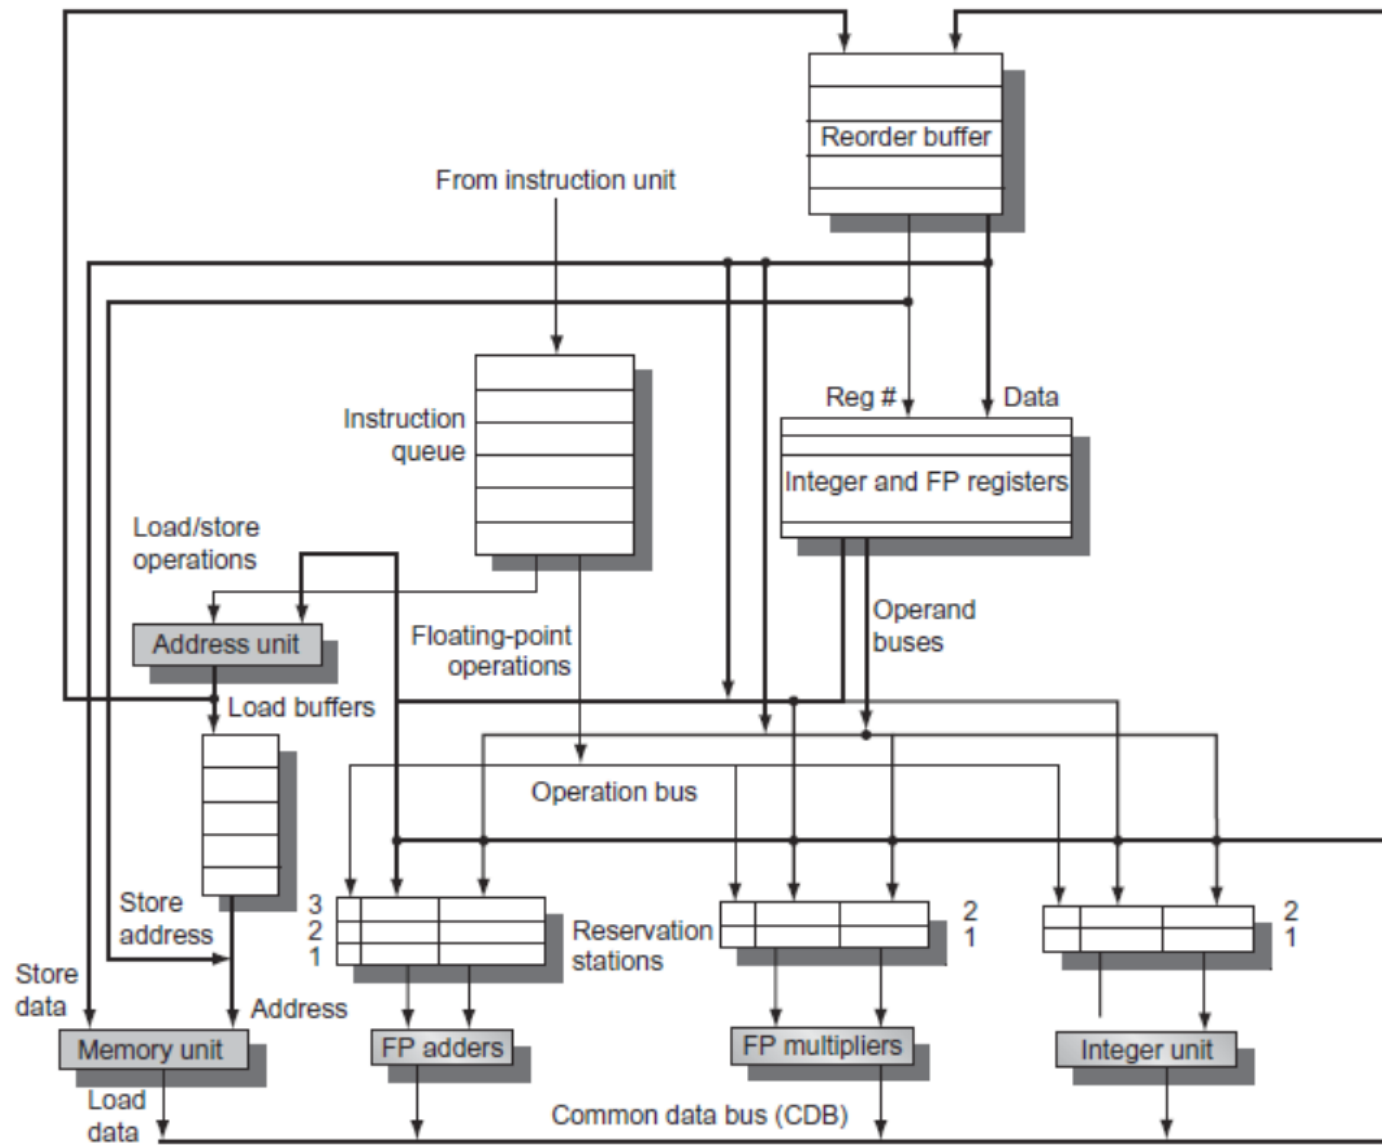
Multiple Issue

- Examine all the dependencies amoung the instructions in the bundle

- If dependencies exist in bundle, encode them in reservation stations

- Also need multiple completion/commit

- To simplify RS allocation:
  - Limit the number of instructions of a given class that can be issued in a "bundle", i.e. on FP, one integer, one load, one store

45

Example

Loop: ld x2,0(x1)   //x2=array element

addi x2,x2,1   //increment x2

sd x2,0(x1)   //store result

addi x1,x1,8   //increment pointer

bne x2,x3,Loop   //branch if not last

Example (No Speculation)

Iteration number	Instructions	Issues at clock cycle number	Executes at clock cycle number	Memory access at clock cycle number	Write CDB at clock cycle number	Comment
1	ld x2,0(x1)	1	2	3	4	First issue
1	addi x2,x2,1	1	5		6	Wait for ld
1	sd x2,0(x1)	2	3	7		Wait for addi
1	addi x1,x1,8	2	3		4	Execute directly
1	bne x2,x3,Loop	3	7			Wait for addi
2	ld x2,0(x1)	4	8	9	10	Wait for bne
2	addi x2,x2,1	4	11		12	Wait for ld
2	sd x2,0(x1)	5	9	13		Wait for addi
2	addi x1,x1,8	5	8		9	Wait for bne
2	bne x2,x3,Loop	6	13			Wait for addi
3	ld x2,0(x1)	7	14	15	16	Wait for bne
3	addi x2,x2,1	7	17		18	Wait for ld
3	sd x2,0(x1)	8	15	19		Wait for addi
3	addi x1,x1,8	8	14		15	Wait for bne
3	bne x2,x3,Loop	9	19			Wait for addi

Example (Mutiple Issue with Speculation)

Iteration number	Instructions	Issues at clock number	Executes at clock number	Read access at clock number	Write CDB at clock number	Commits at clock number	Comment
1	ld x2,0(x1)	1	2	3	4	5	First issue
1	addi x2,x2,1	1	5		6	7	Wait for ld
1	sd x2,0(x1)	2	3			7	Wait for addi
1	addi x1,x1,8	2	3		4	8	Commit in order
1	bne x2,x3,Loop	3	7			8	Wait for addi
2	ld x2,0(x1)	4	5	6	7	9	No execute delay
2	addi x2,x2,1	4	8		9	10	Wait for ld
2	sd x2,0(x1)	5	6			10	Wait for addi
2	addi x1,x1,8	5	6		7	11	Commit in order
2	bne x2,x3,Loop	6	10			11	Wait for addi
3	ld x2,0(x1)	7	8	9	10	12	Earliest possible
3	addi x2,x2,1	7	11		12	13	Wait for ld
3	sd x2,0(x1)	8	9			13	Wait for addi
3	addi x1,x1,8	8	9		10	14	Executes earlier
3	bne x2,x3,Loop	9	13			14	Wait for addi

Dynamic Scheduling, Multiple Issue, and Speculation

Branch-Target Buffer

- Need high instruction bandwidth
  - Branch-Target buffers
    - Next PC prediction buffer, indexed by current PC
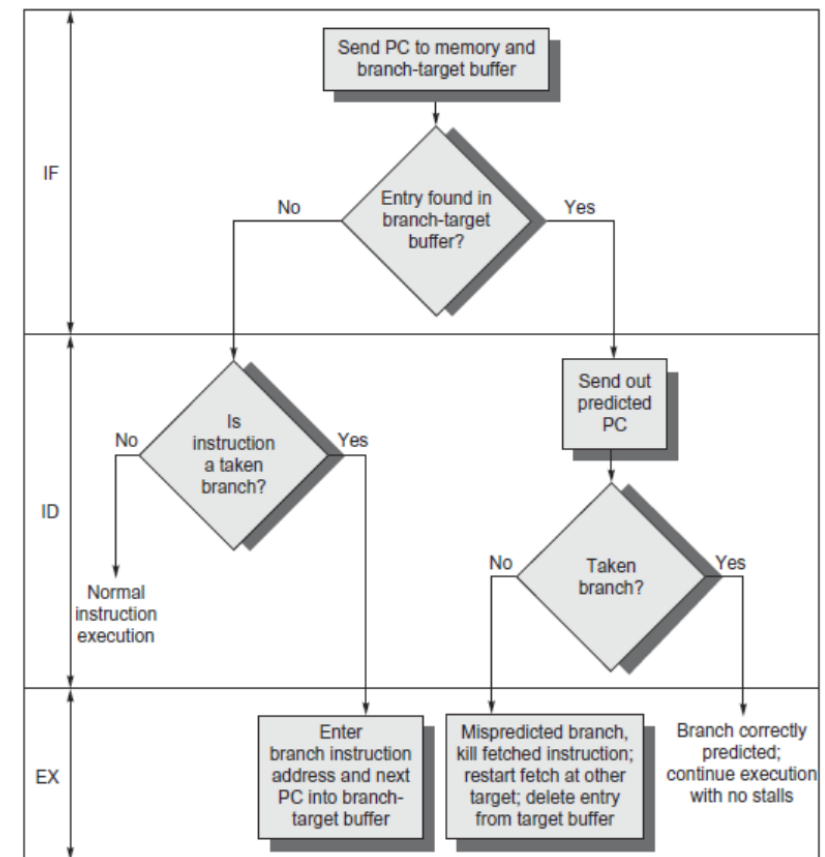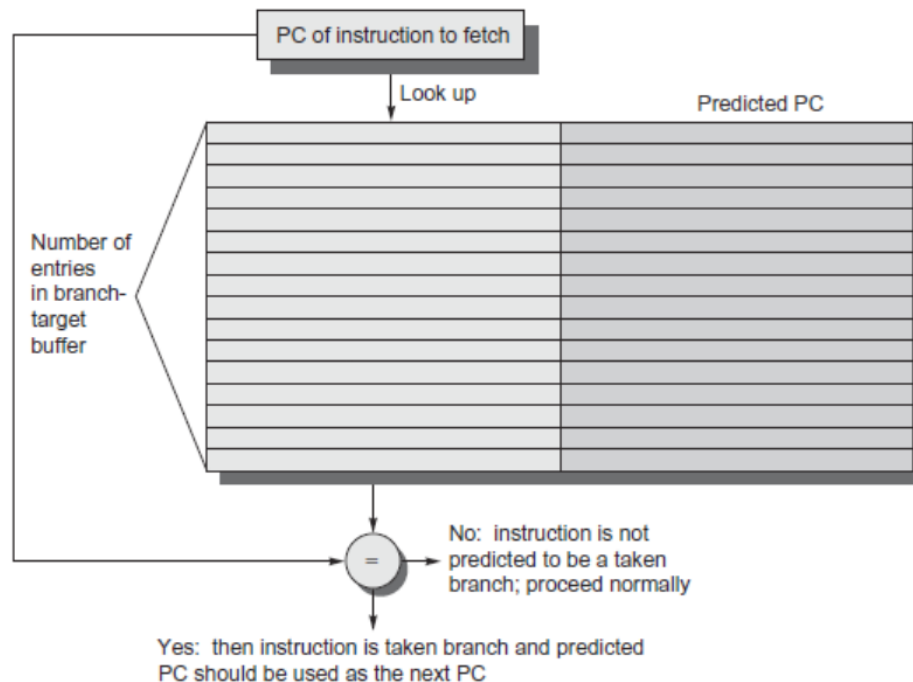
Branch Folding

- Optimization:
  - Larger branch-target buffer
  - Add target instruction into buffer to deal with longer decoding time required by larger buffer
  - "Branch folding"

50

Return Address Predictor

- Most unconditional branches come from function returns

- The same procedure can be called from multiple sites
  - Causes the buffer to potentially forget about the return address from previous calls

- Create return address buffer organized as a stack

51

Return Address Predictor

Integrated Instruction Fetch Unit

- Design monolithic unit that performs:

  - Branch prediction

  - Instruction prefetch

    - Fetch ahead

  - Instruction memory access and buffering

    - Deal with crossing cache lines

Register Renaming

- Register renaming vs. reorder buffers
  - Instead of virtual registers from reservation stations and reorder buffer, create a single register pool
    - Contains visible registers and virtual registers
  - Use hardware-based map to rename registers during issue
  - WAW and WAR hazards are avoided
  - Speculation recovery occurs by copying during commit
  - Still need a ROB-like queue to update table in order
  - Simplifies commit:
    - Record that mapping between architectural register and physical register is no longer speculative
    - Free up physical register used to hold older value
    - In other words: SWAP physical registers on commit
  - Physical register de-allocation is more difficult
    - Simple approach: deallocate virtual register when next instruction writes to its mapped architecturally-visibly register

# Integrated Issue and Renaming

- Combining instruction issue with register renaming:
  - Issue logic pre-reserves enough physical registers for the bundle
  - Issue logic finds dependencies within bundle, maps registers as necessary
  - Issue logic finds dependencies between current bundle and already in-flight bundles, maps registers as necessary

Instr. #	Instruction	Physical register assigned or destination	Instruction with physical register numbers	Rename map changes
1	add x1,x2,x3	p32	add p32,p2,p3	x1-> p32
2	sub x1,x1,x2	p33	sub p33,p32,p2	x1->p33
3	add x2,x1,x2	p34	add p34,p33,x2	x2->p34
4	sub x1,x3,x2	p35	sub p35,p3,p34	x1->p35
5	add x1,x1,x2	p36	add p36,p35,p34	x1->p36
6	sub x1,x3,x1	p37	sub p37,p3,p36	x1->p37

How Much?

- ## How much to speculate

  - ### Mis-speculation degrades performance and power relative to no speculation

    - May cause additional misses (cache, TLB)

  - ### Prevent speculative code from causing higher costing misses (e.g. L2)

- ## Speculating through multiple branches

  - ### Complicates speculation recovery

- ## Speculation and energy efficiency

  - ### Note: speculation is only energy efficient when it significantly improves performance

How Much?

integer

Energy Efficiency

- Value prediction
  - Uses:
    - Loads that load from a constant pool
    - Instruction that produces a value from a small set of values
  - Not incorporated into modern processors
  - Similar idea--*address aliasing prediction*--is used on some processors to determine if two stores or a load and a store reference the same address to allow for reordering

58

Fallacies and Pitfalls

- It is easy to predict the performance/energy efficiency of two different versions of the same ISA if we hold the technology constant

Fallacies and Pitfalls

- Processors with lower CPIs / faster clock rates will also be faster

Processor	Implementation technology	Clock rate	Power	SPECCInt2006 base	SPECCFP2006 baseline
Intel Pentium 4 670	90 nm	3.8 GHz	115 W	11.5	12.2
Intel Itanium 2	90 nm	1.66 GHz	104 W approx. 70 W one core	14.5	17.3
Intel i7 920	45 nm	3.3 GHz	130 W total approx. 80 W one core	35.5	38.4

- Pentium 4 had higher clock, lower CPI
- Itanium had same CPI, lower clock

Fallacies and Pitfalls

- Sometimes bigger and dumber is better
  - Pentium 4 and Itanium were advanced designs, but could not achieve their peak instruction throughput because of relatively small caches as compared to i7

- And sometimes smarter is better than bigger and dumber
  - TAGE branch predictor outperforms gshare with less stored predictions

Fallacies and Pitfalls

- Believing that there are large amounts of ILP available, if only we had the right techniques